

Installation Guide

Latest version: [installation-openstack-ubuntu-note.md](https://github.com/littlewey/installation-openstack-ubuntu-note.md) hosted on *Github.com/littlewey*

ref: <https://docs.openstack.org/install-guide>

Ubuntu was chosen as host OS.

"It's a good way to learn by installing it manually for as many services as you could :-) ."

Wey Gu

Installation Guide

[Security](#)

[Host networking](#)

[My network solution](#)

[Base Machine](#)

[Controller actions](#)

[management network eth0 \(enp0s3\)](#)

[hostname and hosts](#)

[SQL database](#)

[Message queue](#)

[Memcached](#)

[Compute actions](#)

[management network eth0 \(enp0s3\)](#)

[configure NTP by editing `/etc/chrony/chrony.conf`](#)

[Keystone installation](#)

[Configure the Apache HTTP server](#)

[Finalize the installation](#)

[Create a domain, projects, users, and roles](#)

[Verify operation](#)

[Create OpenStack client environment scripts](#)

[Creating the scripts](#)

[Using the scripts](#)

[Glance installation](#)

[Install and configure](#)

[Prerequisites](#)

[Install and configure components](#)

[Finalize installation](#)

[Verify operation](#)

[Nova installation](#)

[Nova install and configure controller node](#)

[Prerequisites](#)

[Install and configure components](#)

[Finalize installation](#)

[Nova Install and configure a compute node](#)

[Install and configure components](#)

[Finalize installation](#)

[Add the compute node to the cell database](#)

[Neutron installation](#)

[Neutron Install and configure controller node](#)

[Prerequisites](#)

[Configure networking options](#)

[Networking Option 1: Provider networks](#)

[Install the components](#)

[Configure the server component](#)

[Configure the Modular Layer 2 \(ML2\) plug-in](#)

[Configure the Linux bridge agent](#)

[Configure the DHCP agent](#)

[Configure the metadata agent](#)

[Configure the Compute service to use the Networking service](#)

[Finalize installation](#)

[Verify operation](#)

[Neutron Install and configure compute node](#)

[Install the components](#)

[Configure the common component](#)

[Configure networking options](#)

[Configure the Linux bridge agent](#)

[Configure the Compute service to use the Networking service](#)

[Finalize installation](#)

[Verify operation](#)

[Congratulations! Let's try booting an instance](#)

[Create provider network/subnetwork](#)

[Create flavor](#)

[Add security group rules](#)

[Launch an instance](#)

[Determine instance options](#)

[Launch the instance](#)

[Access the instance using the virtual console](#)

[Access the instance remotely](#)

[\[ISSUE\] DHCP failure in VM troubleshooting](#)

[what it is like](#)

[Conclusion:](#)

[Cinder](#)

[Cinder on controller](#)

[Install and configure controller node](#)

[Prerequisites](#)

[Install and configure components](#)

[Configure Compute to use Block Storage](#)

[Finalize installation](#)

[Cinder on block storage node](#)

[configure storage network for compute](#)

[Create cinder machine: storage](#)

[Storage actions & cinder on storage node](#)

[Management net eth0 \(enp0s3\) and storage net eth2 \(enp0s9\)](#)

[configure NTP by editing `/etc/chrony/chrony.conf`](#)

[Check new disk was there already](#)

[Install and configure a storage node](#)

[Prerequisites](#)

[Install and configure components](#)

[Finalize installation](#)

[Verify operation](#)

[Let's try something on block storage!](#)

[Create a volume](#)

[Attach the volume to an instance](#)

Security

ref: <https://docs.openstack.org/install-guide/environment-security.html>

```
$ openssl rand -hex 10
```

Host networking

ref: <https://docs.openstack.org/install-guide/environment-networking.html>

ref: <https://help.ubuntu.com/lts/serverguide/network-configuration.html>

The example architectures assume use of the following networks:

- Management on 10.0.0.0/24 with gateway 10.0.0.1
This network requires a gateway to provide Internet access to all nodes for administrative purposes such as package installation, security updates, DNS, and NTP.
- Provider on 203.0.113.0/24 with gateway 203.0.113.1
This network requires a gateway to provide Internet access to instances in your OpenStack environment.

My network solution

Net0:

Network name: VirtualBox host-only Ethernet Adapter
Purpose: administrator / management network
IP block: 10.20.0.0/24
DHCP: disable
Linux device: eth0

Net1:

Network name: VirtualBox host-only Ethernet Adapter#2
Purpose: Provider network
DHCP: disable
IP block: 172.16.0.0/24
Linux device: eth1

Net2:

Network name: VirtualBox host-only Ethernet Adapter#3
Purpose: Storage network
DHCP: disable
IP block: 192.168.199.0/24
Linux device: eth2

Net3:

Network name: VirtualBox Bridged // for accessing network or remote access purpose
Purpose: Internet
DHCP: enable
IP block: <depend on your network>
Linux device: eth3

Edit the `/etc/network/interfaces` file to contain the following:

Replace `INTERFACE_NAME` with the actual interface name. For example, *eth1* or *ens224*.

```
# The provider network interface
auto INTERFACE_NAME
iface INTERFACE_NAME inet manual
up ip link set dev $IFACE up
down ip link set dev $IFACE down
```

Base Machine

- download image from <https://launchpad.net/ubuntu/+mirror/mirrors.neusoft.edu.cn-release>
- Change root password

```
$ sudo su
# passwd
```

- Allow root ssh with password

```
# vi /etc/ssh/sshd_config

PermitRootLogin yes
```

- Check nic names

```
root@ubuntu:~# dmesg | grep rename
[ 2.799294] e1000 0000:00:09.0 enp0s9: renamed from eth2
[ 2.800192] e1000 0000:00:0a.0 enp0s10: renamed from eth3
[ 2.801072] e1000 0000:00:08.0 enp0s8: renamed from eth1
[ 2.804067] e1000 0000:00:03.0 enp0s3: renamed from eth0
```

- configure management network as a dummy one

```
# vi /etc/network/interfaces
auto enp0s3
iface enp0s3 inet static
address 10.20.0.11
netmask 255.255.255.0
```

- NTP

- install chrony

```
install chrony
```

- Edit the `/etc/chrony/chrony.conf` file and add, change, or remove these keys as necessary for your environment:

```
allow 10.20.0.0/24
```

- restart service

```
# service chrony restart
```

- Install OpenStack packages

ref: <https://docs.openstack.org/install-guide/environment-packages.html>

Enable the OpenStack repository

```
# apt install software-properties-common
# add-apt-repository cloud-archive:ocata
```

Upgrade the packages on all nodes:

Set apt proxy before doing that will help save your life

```
# vi /etc/apt/apt.conf.d/90proxy
Acquire::http::Proxy "http://www-proxy.exu.ericsson.se:8080";
Acquire::https::Proxy "http://www-proxy.exu.ericsson.se:8080";

# sed -i -e 's/cn/us/g' /etc/apt/sources.list
```

```
# apt update && apt dist-upgrade -y
```

Install the OpenStack client:

```
# apt install python-openstackclient -y
```

Controller actions

management network eth0 (enp0s3)

```
# vi /etc/network/interfaces

auto enp0s3
iface enp0s3 inet static
address 10.20.0.10
netmask 255.255.255.0

# ifup enp0s3
```

hostname and hosts

```
# echo controller > /etc/hostname
# echo 10.20.0.10    controller >> /etc/hosts
# echo 10.20.0.20    compute >> /etc/hosts
# hostname controller
```

SQL database

Install package

```
# apt install mariadb-server python-pymysql -y
```

Create and edit the `/etc/mysql/mariadb.conf.d/99-openstack.cnf` file and complete the following actions:

Create a `[mysqld]` section, and set the `bind-address` key to the management IP address of the controller node to enable access by other nodes via the management network. Set additional keys to enable useful options and the UTF-8 character set:

```
[mysqld]
bind-address = 10.20.0.10

default-storage-engine = innodb
innodb_file_per_table = on
max_connections = 4096
collation-server = utf8_general_ci
character-set-server = utf8
```

restart database service

```
# service mysql restart
```

Secure the database service by running the `mysql_secure_installation` script. In particular, choose a suitable password for the database `root` account:

```
# mysql_secure_installation
```

Message queue

Install the package:

```
# apt install rabbitmq-server
```

Add the `openstack` user:

```
# rabbitmqctl add_user openstack RABBIT_PASS

Creating user "openstack" ...
```

Replace `RABBIT_PASS` with a suitable password.

Permit configuration, write, and read access for the `openstack` user:

```
# rabbitmqctl set_permissions openstack ".*" ".*" ".*"

Setting permissions for user "openstack" in vhost "/" ...
```

Memcached

Install the packages:

```
# apt install memcached python-memcache
```

Edit the `/etc/memcached.conf` file and configure the service to use the management IP address of the controller node. This is to enable access by other nodes via the management network:

```
-l 10.20.0.10
```

Change the existing line that had `-1 127.0.0.1`.

Restart the Memcached service:

```
# service memcached restart
```

Compute actions

management network eth0 (enp0s3)

```
# vi /etc/network/interfaces

auto enp0s3
iface enp0s3 inet static
address 10.20.0.10
netmask 255.255.255.0

# ifup enp0s3
```

configure NTP by editing `/etc/chrony/chrony.conf`

```
server 10.20.0.10 iburst
```

change hostname and hosts

```
# echo compute > /etc/hostname
# echo 10.20.0.10    controller >> /etc/hosts
# echo 10.20.0.20    compute >> /etc/hosts
# hostname compute
```

Keystone installation

ref: <https://docs.openstack.org/newton/install-guide-ubuntu/keystone.html>

Before you configure the OpenStack Identity service, you must create a database and an administration token.

To create the database, complete the following actions:

- Use the database access client to connect to the database server as the `root` user:

```
$ mysql -u root -p
```

In 16.04 LTS local access need no user/psw

```
# mysql
```


- Create the `keystone` database:

```
mysql> CREATE DATABASE keystone;
```

- Grant proper access to the `keystone` database:

```
mysql> GRANT ALL PRIVILEGES ON keystone.* TO 'keystone'@'localhost' \
IDENTIFIED BY 'KEYSTONE_DBPASS';
mysql> GRANT ALL PRIVILEGES ON keystone.* TO 'keystone'@'%' \
IDENTIFIED BY 'KEYSTONE_DBPASS';
```

Replace `KEYSTONE_DBPASS` with a suitable password.

- Exit the database access client.

Run the following command to install the packages:

```
# apt install keystone -y
```

1. Edit the `/etc/keystone/keystone.conf` file and complete the following actions:

- In the `[database]` section, configure database access:

```
[database]
...
connection = mysql+pymysql://keystone:KEYSTONE_DBPASS@controller/keystone
```

Replace `KEYSTONE_DBPASS` with the password you chose for the database.

Comment out or remove any other `connection` options in the `[database]` section.

- In the `[token]` section, configure the Fernet token provider:

```
[token]
...
provider = fernet
```

2. Populate the Identity service database:

```
# su -s /bin/sh -c "keystone-manage db_sync" keystone
```

3. Initialize Fernet key repositories:

```
# keystone-manage fernet_setup --keystone-user keystone --keystone-group keystone
# keystone-manage credential_setup --keystone-user keystone --keystone-group keystone
```

4. Bootstrap the Identity service:

```
# keystone-manage bootstrap --bootstrap-password ADMIN_PASS \  
--bootstrap-admin-url http://controller:35357/v3/ \  
--bootstrap-internal-url http://controller:35357/v3/ \  
--bootstrap-public-url http://controller:5000/v3/ \  
--bootstrap-region-id RegionOne
```

Replace `ADMIN_PASS` with a suitable password for an administrative user.

Configure the Apache HTTP server

1. Edit the `/etc/apache2/apache2.conf` file and configure the `ServerName` option to reference the controller node:

```
ServerName controller
```

Finalize the installation

1. Restart the Apache service and remove the default SQLite database:

```
# service apache2 restart  
# rm -f /var/lib/keystone/keystone.db
```

1. Configure the administrative account

```
$ export OS_USERNAME=admin  
$ export OS_PASSWORD=ADMIN_PASS  
$ export OS_PROJECT_NAME=admin  
$ export OS_USER_DOMAIN_NAME=Default  
$ export OS_PROJECT_DOMAIN_NAME=Default  
$ export OS_AUTH_URL=http://controller:35357/v3  
$ export OS_IDENTITY_API_VERSION=3
```

Replace `ADMIN_PASS` with the password used in the `keystone-manage bootstrap` command from the section called [Install and configure](#).

Create a domain, projects, users, and roles

The Identity service provides authentication services for each OpenStack service. The authentication service uses a combination of [domains](#), [projects](#), [users](#), and [roles](#).

1. This guide uses a service project that contains a unique user for each service that you add to your environment. Create the `service` project:

```
$ openstack project create --domain default \
  --description "Service Project" service
```

Field	Value
description	Service Project
domain_id	default
enabled	True
id	24ac7f19cd944f4cba1d77469b2a73ed
is_domain	False
name	service
parent_id	default

2. Regular (non-admin) tasks should use an unprivileged project and user. As an example, this guide creates the `demo` project and user.

- Create the `demo` project:

```
$ openstack project create --domain default \
  --description "Demo Project" demo
```

Field	Value
description	Demo Project
domain_id	default
enabled	True
id	231ad6e7ebba47d6a1e57e1cc07ae446
is_domain	False
name	demo
parent_id	default

Do not repeat this step when creating additional users for this project.

- Create the `demo` user:

```
$ openstack user create --domain default \
  --password-prompt demo
```

User Password:

Repeat User Password:

Field	Value
domain_id	default
enabled	True
id	aeda23aa78f44e859900e22c24817832
name	demo
password_expires_at	None

- Create the `user` role:

```
$ openstack role create user
```

Field	Value
domain_id	None
id	997ce8d05fc143ac97d83fdfb5998552
name	user

- Add the `user` role to the `demo` project and user:

```
$ openstack role add --project demo --user demo user
```

Verify operation

For security reasons, disable the temporary authentication token mechanism:

Edit the `/etc/keystone/keystone-paste.ini` file and remove `admin_token_auth` from the `[pipeline:public_api]`, `[pipeline:admin_api]`, and `[pipeline:api_v3]` sections.

Unset the temporary `OS_AUTH_URL` and `OS_PASSWORD` environment variable:

```
$ unset OS_AUTH_URL OS_PASSWORD
```

As the `admin` user, request an authentication token:

```
$ openstack --os-auth-url http://controller:35357/v3 \
--os-project-domain-name Default --os-user-domain-name Default \
--os-project-name admin --os-username admin token issue
```

Password:

Field	Value
expires	2017-08-12T20:14:07.056119Z
id	gAAAAABWvi7_B8kKQD9wdXac8MoZiQldmjE0643d-e_j-XXq9AmIegIbA7UHGpV
	atnN21qtOMjCFWX7BREJEQnVOAj3nc1RQgAYRsfsU_MrsuWb4EDtnjU7HEpoBb4
	o6ozsA_NmFWEpLeKy0uNn_WeKbAhYygrsmQGA49dc1HVnz-OMVLIyM9ws
project_id	343d245e850143a096806dfaefa9afdc
user_id	ac3377633149401296f6c0d92d79dc16

This command uses the password for the `admin` user. As we gave above it's `ADMIN_PASS`.

As the `demo` user, request an authentication token:

```
$ openstack --os-auth-url http://controller:5000/v3 \
--os-project-domain-name Default --os-user-domain-name Default \
--os-project-name demo --os-username demo token issue
```

Password:

Field	Value
expires	2017-08-12T20:15:39.014479Z
id	gAAAAABWvi9bsh7vkiBy5BpCCnc-JkbGhm9wH3fabS_cY7uab0ubesi-Me6IGWW
	yQqNegDDZ5jw7grI26vvgy1J5nCVwZ_zFRqPiz_qhbq29mgbQLglbkq6FQvzBRQ
	JcOzq3uwhzNxsZJWmzGC7rJE_H0A_a3UFhqv8M4zMRYSbS2YF0MyFmp_U
project_id	ed0b60bf607743088218b0a533d5943f
user_id	58126687cbcc4888bfa9ab73a2256f27

This command uses the password for the `demo` user and API port 5000 which only allows regular (non-admin) access to the Identity service API.

Create OpenStack client environment scripts

The previous section used a combination of environment variables and command options to interact with the Identity service via the `openstack` client. To increase efficiency of client operations, OpenStack supports simple client environment scripts also known as OpenRC files. These scripts typically contain common options for all clients, but also support unique options. For more information, see the [OpenStack End User Guide](#).

Creating the scripts

Create client environment scripts for the `admin` and `demo` projects and users. Future portions of this guide reference these scripts to load appropriate credentials for client operations.

1. Edit the `admin-openrc` file and add the following content:

```
export OS_PROJECT_DOMAIN_NAME=Default
export OS_USER_DOMAIN_NAME=Default
export OS_PROJECT_NAME=admin
export OS_USERNAME=admin
export OS_PASSWORD=ADMIN_PASS
export OS_AUTH_URL=http://controller:35357/v3
export OS_IDENTITY_API_VERSION=3
export OS_IMAGE_API_VERSION=2
```

Replace `ADMIN_PASS` with the password you chose for the `admin` user in the Identity service.

2. Edit the `demo-openrc` file and add the following content:

```
export OS_PROJECT_DOMAIN_NAME=Default
export OS_USER_DOMAIN_NAME=Default
export OS_PROJECT_NAME=demo
export OS_USERNAME=demo
export OS_PASSWORD=demo
export OS_AUTH_URL=http://controller:5000/v3
export OS_IDENTITY_API_VERSION=3
export OS_IMAGE_API_VERSION=2
```

Replace `OS_PASSWORD=demo` with the password you chose for the `demo` user in the Identity service.

Using the scripts

To run clients as a specific project and user, you can simply load the associated client environment script prior to running them. For example:

1. Load the `admin-openrc` file to populate environment variables with the location of the Identity service and the `admin` project and user credentials:

```
$ . admin-openrc
```

2. Request an authentication token:

```
root@ubuntu:~# openstack token issue --max-width 70
+-----+-----+
| Field      | Value                                                                 |
+-----+-----+
| expires    | 2017-08-23T14:00:10+0000                                             |
| id         | gAAAAABZnXxaKuuwh9Kw-dbY1mSn8LeNNQMkMIj2EW8jyj00NSy5H             |
|            | QPno4Tj6NqqSkumKhRZW81PS1nZC2pm5fCuH5XMtVfJTU89RX6Sba             |
|            | -vSv-OZ15uHvRY4KOK03WH15Dnp1XbWN97xY8tR_kAhc-69                   |
|            | -WvDe1DLS6vKr-bKbYDVXLq1LshE8E                                     |
| project_id | 78c9c849237649a3a8c4526167427589                                     |
| user_id    | d8efd16c30904a7992010abe4bdb9a2b                                     |
+-----+-----+
```

Glance installation

ref: <https://docs.openstack.org/newton/install-guide-ubuntu/glance.html>

For simplicity, this guide describes configuring the Image service to use the `file` back end, which uploads and stores in a directory on the controller node hosting the Image service. By default, this directory is

```
/var/lib/glance/images/.
```

Before you proceed, ensure that the controller node has at least several gigabytes of space available in this directory. Keep in mind that since the `file` back end is often local to a controller node, it is not typically suitable for a multi-node glance deployment.

For information on requirements for other back ends, see [Configuration Reference](#).

Install and configure

This section describes how to install and configure the Image service, code-named glance, on the controller node. For simplicity, this configuration stores images on the local file system.

Prerequisites

Before you install and configure the Image service, you must create a database, service credentials, and API endpoints.

1. To create the database, complete these steps:

- Use the database access client to connect to the database server as the `root` user:

```
$ mysql -u root -p
```

- Create the `glance` database:

```
mysql> CREATE DATABASE glance;
```

- Grant proper access to the `glance` database:

```
mysql> GRANT ALL PRIVILEGES ON glance.* TO 'glance'@'localhost' \
IDENTIFIED BY 'GLANCE_DBPASS';
mysql> GRANT ALL PRIVILEGES ON glance.* TO 'glance'@'%' \
IDENTIFIED BY 'GLANCE_DBPASS';
```

Replace `GLANCE_DBPASS` with a suitable password.

- Exit the database access client.

2. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

3. To create the service credentials, complete these steps:

- Create the `glance` user:

```
$ openstack user create --domain default --password-prompt glance
```

User Password:

Repeat User Password:

Field	Value
domain_id	default
enabled	True
id	3f4e777c4062483ab8d9edd7dffa829df
name	glance
password_expires_at	None

- Add the `admin` role to the `glance` user and `service` project:

```
$ openstack role add --project service --user glance admin
```

This command provides no output.

- Create the `glance` service entity:

```
$ openstack service create --name glance \
  --description "OpenStack Image" image
```

Field	Value
description	OpenStack Image
enabled	True
id	8c2c7f1b9b5049ea9e63757b5533e6d2
name	glance
type	image

4. Create the Image service API endpoints:


```
$ openstack endpoint create --region RegionOne \  
  image public http://controller:9292
```

```
+-----+-----+  
| Field      | Value                                |  
+-----+-----+  
| enabled    | True                                |  
| id         | 340be3625e9b4239a6415d034e98aace |  
| interface  | public                              |  
| region     | RegionOne                           |  
| region_id  | RegionOne                           |  
| service_id | 8c2c7f1b9b5049ea9e63757b5533e6d2 |  
| service_name | glance                              |  
| service_type | image                               |  
| url        | http://controller:9292             |  
+-----+-----+
```

```
$ openstack endpoint create --region RegionOne \  
  image internal http://controller:9292
```

```
+-----+-----+  
| Field      | Value                                |  
+-----+-----+  
| enabled    | True                                |  
| id         | a6e4b153c2ae4c919eccfdbb7dceb5d2 |  
| interface  | internal                            |  
| region     | RegionOne                           |  
| region_id  | RegionOne                           |  
| service_id | 8c2c7f1b9b5049ea9e63757b5533e6d2 |  
| service_name | glance                              |  
| service_type | image                               |  
| url        | http://controller:9292             |  
+-----+-----+
```

```
$ openstack endpoint create --region RegionOne \  
  image admin http://controller:9292
```

```
+-----+-----+  
| Field      | Value                                |  
+-----+-----+  
| enabled    | True                                |  
| id         | 0c37ed58103f4300a84ff125a539032d |  
| interface  | admin                              |  
| region     | RegionOne                           |  
| region_id  | RegionOne                           |  
| service_id | 8c2c7f1b9b5049ea9e63757b5533e6d2 |  
| service_name | glance                              |  
| service_type | image                               |  
| url        | http://controller:9292             |  
+-----+-----+
```

Install and configure components

Install the packages:

```
# apt install glance -y
```

1. Edit the `/etc/glance/glance-api.conf` file and complete the following actions:

- In the `[database]` section, configure database access:

```
[database]
...
connection = mysql+pymysql://glance:GLANCE_DBPASS@controller/glance
```

Replace `GLANCE_DBPASS` with the password you chose for the Image service database.

- In the `[keystone_authtoken]` and `[paste_deploy]` sections, configure Identity service access:

```
[keystone_authtoken]
...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = Default
user_domain_name = Default
project_name = service
username = glance
password = glance

[paste_deploy]
...
flavor = keystone
```

Replace `password = glance` with the password you chose for the `glance` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

- In the `[glance_store]` section, configure the local file system store and location of image files:

```
[glance_store]
...
stores = file,http
default_store = file
filesystem_store_datadir = /var/lib/glance/images/
```

2. Edit the `/etc/glance/glance-registry.conf` file and complete the following actions:

- In the `[database]` section, configure database access:

```
[database]
...
connection = mysql+pymysql://glance:GLANCE_DBPASS@controller/glance
```

Replace `GLANCE_DBPASS` with the password you chose for the Image service database.

- In the `[keystone_authtoken]` and `[paste_deploy]` sections, configure Identity service access:

```
[keystone_authtoken]
...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = Default
user_domain_name = Default
project_name = service
username = glance
password = glance

[paste_deploy]
...
flavor = keystone
```

Replace `password = glance` with the password you chose for the `glance` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

Populate the Image service database:

```
# su -s /bin/sh -c "glance-manage db_sync" glance
```

Ignore any deprecation messages in this output.

Finalize installation

Restart the Image services:

```
# service glance-registry restart
# service glance-api restart
```

Verify operation

Verify operation of the Image service using [CirrOS](#), a small Linux image that helps you test your OpenStack deployment.

For more information about how to download and build images, see [OpenStack Virtual Machine Image Guide](#). For information about how to manage images, see the [OpenStack End User Guide](#).

1. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

2. Download the source image:

```
$ wget http://download.cirros-cloud.net/0.3.4/cirros-0.3.4-x86_64-disk.img
```

tip: add proxy to improve speed in office network

```
$ export http_proxy=http://www-proxy.exu.ericsson.se:8080

// after wget

$ unset http_proxy
```

Install `wget` if your distribution does not include it.

3. Upload the image to the Image service using the [QCOW2](#) disk format, [bare](#) container format, and public visibility so all projects can access it:

```
$ openstack image create "cirros" \
  --file cirros-0.3.4-x86_64-disk.img \
  --disk-format qcow2 --container-format bare \
  --public
```

```
+-----+-----+
| Field          | Value                                                                 |
+-----+-----+
| checksum       | 133eae9fb1c98f45894a4e60d8736619                                     |
| container_format | bare                                                                    |
| created_at     | 2015-03-26T16:52:10Z                                                 |
| disk_format    | qcow2                                                                  |
| file           | /v2/images/cc5c6982-4910-471e-b864-1098015901b5/file                 |
| id             | cc5c6982-4910-471e-b864-1098015901b5                                 |
| min_disk       | 0                                                                       |
| min_ram        | 0                                                                       |
| name           | cirros                                                                |
| owner          | ae7a98326b9c45558edd2656d723b9d                                     |
| protected      | False                                                                  |
| schema         | /v2/schemas/image                                                    |
| size           | 13200896                                                              |
| status         | active                                                                |
| tags           |                                                                         |
| updated_at     | 2015-03-26T16:52:10Z                                                 |
| virtual_size   | None                                                                  |
| visibility     | public                                                                |
+-----+-----+
```

For information about the **openstack image create** parameters, see [Create or update an image \(glance\)](#) in the `OpenStack UserGuide`.

For information about disk and container formats for images, see [Disk and container formats for images](#) in the `OpenStack VirtualMachine Image Guide`.

OpenStack generates IDs dynamically, so you will see different values in the example command output.

4. Confirm upload of the image and validate attributes:

```
$ openstack image list
```

```
+-----+-----+-----+
| ID                               | Name   | Status |
+-----+-----+-----+
| 38047887-61a7-41ea-9b49-27987d5e8bb9 | cirros | active |
+-----+-----+-----+
```

Nova installation

ref: <https://docs.openstack.org/newton/install-guide-ubuntu/nova.html>

Nova install and configure controller node

Prerequisites

Before you install and configure the Compute service, you must create databases, service credentials, and API endpoints.

1. To create the databases, complete these steps:

- Use the database access client to connect to the database server as the `root` user:

```
# mysql
```

- Create the `nova_api`, `nova`, and `nova_cell0` databases:

```
MariaDB [(none)]> CREATE DATABASE nova_api;
MariaDB [(none)]> CREATE DATABASE nova;
MariaDB [(none)]> CREATE DATABASE nova_cell0;
```

- Grant proper access to the databases:

```

MariaDB [(none)]> GRANT ALL PRIVILEGES ON nova_api.* TO 'nova'@'localhost' \
IDENTIFIED BY 'NOVA_DBPASS';
MariaDB [(none)]> GRANT ALL PRIVILEGES ON nova_api.* TO 'nova'@'%' \
IDENTIFIED BY 'NOVA_DBPASS';

MariaDB [(none)]> GRANT ALL PRIVILEGES ON nova.* TO 'nova'@'localhost' \
IDENTIFIED BY 'NOVA_DBPASS';
MariaDB [(none)]> GRANT ALL PRIVILEGES ON nova.* TO 'nova'@'%' \
IDENTIFIED BY 'NOVA_DBPASS';

MariaDB [(none)]> GRANT ALL PRIVILEGES ON nova_cell0.* TO 'nova'@'localhost' \
IDENTIFIED BY 'NOVA_DBPASS';
MariaDB [(none)]> GRANT ALL PRIVILEGES ON nova_cell0.* TO 'nova'@'%' \
IDENTIFIED BY 'NOVA_DBPASS';

```

Replace `NOVA_DBPASS` with a suitable password.

- Exit the database access client.

2. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

3. Create the Compute service credentials:

- Create the `nova` user:

```
$ openstack user create --domain default --password-prompt nova
```

User Password:

Repeat User Password:

Field	Value
domain_id	default
enabled	True
id	8a7dbf5279404537b1c7b86c033620fe
name	nova
options	{}
password_expires_at	None

- Add the `admin` role to the `nova` user:

```
$ openstack role add --project service --user nova admin
```

This command provides no output.

- Create the `nova` service entity:

```
$ openstack service create --name nova \  
  --description "OpenStack Compute" compute
```

```
+-----+-----+  
| Field      | Value                                     |  
+-----+-----+  
| description | OpenStack Compute                       |  
| enabled     | True                                    |  
| id          | 060d59eac51b4594815603d75a00aba2      |  
| name        | nova                                    |  
| type        | compute                                |  
+-----+-----+
```

4. Create the Compute API service endpoints:

```
$ openstack endpoint create --region RegionOne \  
compute public http://controller:8774/v2.1
```

Field	Value
enabled	True
id	3c1caa473bfe4390a11e7177894bcc7b
interface	public
region	RegionOne
region_id	RegionOne
service_id	060d59eac51b4594815603d75a00aba2
service_name	nova
service_type	compute
url	http://controller:8774/v2.1

```
$ openstack endpoint create --region RegionOne \  
compute internal http://controller:8774/v2.1
```

Field	Value
enabled	True
id	e3c918de680746a586eac1f2d9bc10ab
interface	internal
region	RegionOne
region_id	RegionOne
service_id	060d59eac51b4594815603d75a00aba2
service_name	nova
service_type	compute
url	http://controller:8774/v2.1

```
$ openstack endpoint create --region RegionOne \  
compute admin http://controller:8774/v2.1
```

Field	Value
enabled	True
id	38f7af91666a47cfb97b4dc790b94424
interface	admin
region	RegionOne
region_id	RegionOne
service_id	060d59eac51b4594815603d75a00aba2
service_name	nova
service_type	compute
url	http://controller:8774/v2.1

5. Create a Placement service user using your chosen `PLACEMENT_PASS`:


```
$ openstack user create --domain default --password-prompt placement
```

User Password:

Repeat User Password:

Field	Value
domain_id	default
enabled	True
id	fa742015a6494a949f67629884fc7ec8
name	placement
options	{}
password_expires_at	None

6. Add the Placement user to the service project with the admin role:

```
$ openstack role add --project service --user placement admin
```

This command provides no output.

7. Create the Placement API entry in the service catalog:

```
$ openstack service create --name placement --description "Placement API" placement
```

Field	Value
description	Placement API
enabled	True
id	2d1a27022e6e4185b86adac4444c495f
name	placement
type	placement

8. Create the Placement API service endpoints:

```
$ openstack endpoint create --region RegionOne placement public http://controller:8778
```

```
+-----+-----+
| Field      | Value                                     |
+-----+-----+
| enabled    | True                                    |
| id         | 2b1b2637908b4137a9c2e0470487cbc0      |
| interface  | public                                 |
| region     | RegionOne                             |
| region_id  | RegionOne                             |
| service_id | 2d1a27022e6e4185b86adac4444c495f     |
| service_name | placement                             |
| service_type | placement                             |
| url        | http://controller:8778                 |
+-----+-----+
```

```
$ openstack endpoint create --region RegionOne placement internal http://controller:8778
```

```
+-----+-----+
| Field      | Value                                     |
+-----+-----+
| enabled    | True                                    |
| id         | 02bcda9a150a4bd7993ff4879df971ab     |
| interface  | internal                               |
| region     | RegionOne                             |
| region_id  | RegionOne                             |
| service_id | 2d1a27022e6e4185b86adac4444c495f     |
| service_name | placement                             |
| service_type | placement                             |
| url        | http://controller:8778                 |
+-----+-----+
```

```
$ openstack endpoint create --region RegionOne placement admin http://controller:8778
```

```
+-----+-----+
| Field      | Value                                     |
+-----+-----+
| enabled    | True                                    |
| id         | 3d71177b9e0f406f98cbff198d74b182     |
| interface  | admin                                 |
| region     | RegionOne                             |
| region_id  | RegionOne                             |
| service_id | 2d1a27022e6e4185b86adac4444c495f     |
| service_name | placement                             |
| service_type | placement                             |
| url        | http://controller:8778                 |
+-----+-----+
```

Install and configure components

Default configuration files vary by distribution. You might need to add these sections and options rather than modifying existing sections and options. Also, an ellipsis (...) in the configuration snippets indicates potential default configuration options that you should retain.

1. Install the packages:

```
# apt install nova-api nova-conductor nova-consoleauth \
nova-novncproxy nova-scheduler nova-placement-api
```

1. Edit the `/etc/nova/nova.conf` file and complete the following actions:

- In the `[api_database]` and `[database]` sections, configure database access:

```
[api_database]
# ...
connection = mysql+pymysql://nova:NOVA_DBPASS@controller/nova_api

[database]
# ...
connection = mysql+pymysql://nova:NOVA_DBPASS@controller/nova
```

Replace `NOVA_DBPASS` with the password you chose for the Compute databases.

- In the `[DEFAULT]` section, configure `RabbitMQ` message queue access:

```
[DEFAULT]
# ...
transport_url = rabbit://openstack:RABBIT_PASS@controller
```

Replace `RABBIT_PASS` with the password you chose for the `openstack` account in `RabbitMQ`.

- In the `[api]` and `[keystone_authtoken]` sections, configure Identity service access:

```
[api]
# ...
auth_strategy = keystone

[keystone_authtoken]
# ...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = default
user_domain_name = default
project_name = service
username = nova
password = nova
```

Replace `nova` with the password you chose for the `nova` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

- In the `[DEFAULT]` section, configure the `my_ip` option to use the management interface IP address of the controller node:

```
[DEFAULT]
# ...
my_ip = 10.0.0.11
```

- In the `[DEFAULT]` section, enable support for the Networking service:

```
[DEFAULT]
# ...
use_neutron = True
firewall_driver = nova.virt.firewall.NoopFirewallDriver
```

By default, Compute uses an internal firewall driver. Since the Networking service includes a firewall driver, you must disable the Compute firewall driver by using the `nova.virt.firewall.NoopFirewallDriver` firewall driver.

- In the `[vnc]` section, configure the VNC proxy to use the management interface IP address of the controller node:

```
[vnc]
enabled = true
# ...
vncserver_listen = $my_ip
vncserver_proxyclient_address = $my_ip
```

- In the `[glance]` section, configure the location of the Image service API:

```
[glance]
# ...
api_servers = http://controller:9292
```

- In the `[oslo_concurrency]` section, configure the lock path:

```
[oslo_concurrency]
# ...
lock_path = /var/lib/nova/tmp
```

- Due to a packaging bug, remove the `log_dir` option from the `[DEFAULT]` section.
- In the `[placement]` section, configure the Placement API:

```
[placement]
# ...
os_region_name = RegionOne
project_domain_name = Default
project_name = service
auth_type = password
user_domain_name = Default
auth_url = http://controller:35357/v3
username = placement
password = PLACEMENT_PASS
```

Replace `PLACEMENT_PASS` with the password you choose for the `placement` user in the Identity service. Comment out any other options in the `[placement]` section.

1. Populate the nova-api database:

```
# su -s /bin/sh -c "nova-manage api_db sync" nova
```

Ignore any deprecation messages in this output.

2. Register the `cell0` database:

```
# su -s /bin/sh -c "nova-manage cell_v2 map_cell0" nova
```

3. Create the `cell1` cell:

```
# su -s /bin/sh -c "nova-manage cell_v2 create_cell --name=cell1 --verbose" nova
109e1d4b-536a-40d0-83c6-5f121b82b650
```

4. Populate the nova database:

```
# su -s /bin/sh -c "nova-manage db sync" nova
```

5. Verify nova cell0 and cell1 are registered correctly:

```
# nova-manage cell_v2 list_cells
+-----+-----+
| Name | UUID |
+-----+-----+
| cell1 | 109e1d4b-536a-40d0-83c6-5f121b82b650 |
| cell0 | 00000000-0000-0000-0000-000000000000 |
+-----+-----+
```

Finalize installation

- Restart the Compute services:

```
# service nova-api restart
# service nova-consoleauth restart
# service nova-scheduler restart
# service nova-conductor restart
# service nova-novncproxy restart
```

Nova Install and configure a compute node

This section describes how to install and configure the Compute service on a compute node. The service supports several [hypervisors](#) to deploy [instances](#) or [VMs](#). For simplicity, this configuration uses the [QEMU](#) hypervisor with the [KVM](#) extension on compute nodes that support hardware acceleration for virtual machines.

Install and configure components

1. Install the packages:

```
# apt install nova-compute
```

1. Edit the `/etc/nova/nova.conf` file and complete the following actions:

- In the `[DEFAULT]` section, configure `RabbitMQ` message queue access:

```
[DEFAULT]
...
transport_url = rabbit://openstack:RABBIT_PASS@controller
```

Replace `RABBIT_PASS` with the password you chose for the `openstack` account in `RabbitMQ`.

- In the `[DEFAULT]` and `[keystone_authtoken]` sections, configure Identity service access:

```
[DEFAULT]
...
auth_strategy = keystone

[keystone_authtoken]
...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = Default
user_domain_name = Default
project_name = service
username = nova
password = nova
```

Replace `password = nova` with the password you chose for the `nova` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

- In the `[DEFAULT]` section, configure the `my_ip` option:

```
[DEFAULT]
...
my_ip = MANAGEMENT_INTERFACE_IP_ADDRESS
```

Replace `MANAGEMENT_INTERFACE_IP_ADDRESS` with the IP address of the management network interface on your compute node, typically 10.0.0.31 for the first node in the [example architecture](#).

here our compute is 10.20.0.20

- In the `[DEFAULT]` section, enable support for the Networking service:

```
[DEFAULT]
...
use_neutron = True
firewall_driver = nova.virt.firewall.NoopFirewallDriver
```

By default, Compute uses an internal firewall service. Since Networking includes a firewall service, you must disable the Compute firewall service by using the `nova.virt.firewall.NoopFirewallDriver` firewall driver.

- In the `[vnc]` section, enable and configure remote console access:

```
[vnc]
...
enabled = True
vncserver_listen = 0.0.0.0
vncserver_proxyclient_address = $my_ip
novncproxy_base_url = http://controller:6080/vnc_auto.html
```

The server component listens on all IP addresses and the proxy component only listens on the management interface IP address of the compute node.

The base URL indicates the location where you can use a web browser to access remote consoles of instances on this compute node.

If the web browser to access remote consoles resides on a host that cannot resolve the `controller` hostname, you must replace `controller` with the management interface IP address of the controller node.

- In the `[glance]` section, configure the location of the Image service API:

```
[glance]
...
api_servers = http://controller:9292
```

- In the `[oslo_concurrency]` section, configure the lock path:

```
[oslo_concurrency]
...
lock_path = /var/lib/nova/tmp
```

- Due to a packaging bug, remove the `log-dir` option from the `[DEFAULT]` section.
- In the `[placement]` section, configure the Placement API:

```
[placement]
# ...
os_region_name = RegionOne
project_domain_name = Default
project_name = service
auth_type = password
user_domain_name = Default
auth_url = http://controller:35357/v3
username = placement
password = placement
```

Replace `placement` with the password you choose for the `placement` user in the Identity service. Comment out any other options in the `[placement]` section.

Finalize installation

Determine whether your compute node supports hardware acceleration for virtual machines:

```
$ egrep -c '(vmx|svm)' /proc/cpuinfo
```

If this command returns a value of `one or greater`, your compute node supports hardware acceleration which typically requires no additional configuration.

If this command returns a value of `zero`, your compute node does not support hardware acceleration and you must configure `libvirt` to use QEMU instead of KVM.

- Edit the `[libvirt]` section in the `/etc/nova/nova-compute.conf` file as follows:

```
[libvirt]
...
virt_type = qemu
```

Restart the Compute service:

```
# service nova-compute restart
```

Add the compute node to the cell database

Run the following commands on the **controller** node.

1. Source the admin credentials to enable admin-only CLI commands, then confirm there are compute hosts in the database:


```
$ . admin-openrc

$ openstack hypervisor list
```

ID	Hypervisor Hostname	Hypervisor Type	Host IP	State
1	compute1	QEMU	10.0.0.31	up

2. Discover compute hosts:

```
# su -s /bin/sh -c "nova-manage cell_v2 discover_hosts --verbose" nova

Found 2 cell mappings.
Skipping cell0 since it does not contain hosts.
Getting compute nodes from cell 'cell1': ad5a5985-a719-4567-98d8-8d148aaaae4bc
Found 1 computes in cell: ad5a5985-a719-4567-98d8-8d148aaaae4bc
Checking host mapping for compute host 'compute': fe58ddc1-1d65-4f87-9456-bc040dc106b3
Creating host mapping for compute host 'compute': fe58ddc1-1d65-4f87-9456-bc040dc106b3
```

When you add new compute nodes, you must run `nova-manage cell_v2 discover_hosts` on the controller node to register those new compute nodes. Alternatively, you can set an appropriate interval in `/etc/nova/nova.conf`:

```
[scheduler]
discover_hosts_in_cells_interval = 300
```

Neutron installation

ref: <https://docs.openstack.org/newton/install-guide-ubuntu/neutron.html>

This chapter explains how to install and configure the Networking service (neutron) using the [provider networks](#).

For more information about the Networking service including virtual networking components, layout, and traffic flows, see the [OpenStack Networking Guide](#).

Neutron Install and configure controller node

Prerequisites

Before you configure the OpenStack Networking (neutron) service, you must create a database, service credentials, and API endpoints.

1. To create the database, complete these steps:

- Use the database access client to connect to the database server as the `root` user:

```
$ mysql -u root -p
```

- Create the `neutron` database:

```
mysql> CREATE DATABASE neutron;
```

- Grant proper access to the `neutron` database, replacing `NEUTRON_DBPASS` with a suitable password:

```
mysql> GRANT ALL PRIVILEGES ON neutron.* TO 'neutron'@'localhost' \
IDENTIFIED BY 'NEUTRON_DBPASS';
mysql> GRANT ALL PRIVILEGES ON neutron.* TO 'neutron'@'%' \
IDENTIFIED BY 'NEUTRON_DBPASS';
```

- Exit the database access client.

2. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

3. To create the service credentials, complete these steps:

- Create the `neutron` user:

```
$ openstack user create --domain default --password-prompt neutron
```

User Password:

Repeat User Password:

Field	Value
domain_id	default
enabled	True
id	319f34694728440eb8ffcb27b6dd8b8a
name	neutron
password_expires_at	None

- Add the `admin` role to the `neutron` user:

```
$ openstack role add --project service --user neutron admin
```

This command provides no output.

- Create the `neutron` service entity:

```
$ openstack service create --name neutron \
  --description "OpenStack Networking" network
```

Field	Value
description	OpenStack Networking
enabled	True
id	f71529314dab4a4d8eca427e701d209e
name	neutron
type	network

4. Create the Networking service API endpoints:

```
$ openstack endpoint create --region RegionOne \  
network public http://controller:9696
```

Field	Value
enabled	True
id	85d80a6d02fc4b7683f611d7fc1493a3
interface	public
region	RegionOne
region_id	RegionOne
service_id	f71529314dab4a4d8eca427e701d209e
service_name	neutron
service_type	network
url	http://controller:9696

```
$ openstack endpoint create --region RegionOne \  
network internal http://controller:9696
```

Field	Value
enabled	True
id	09753b537ac74422a68d2d791cf3714f
interface	internal
region	RegionOne
region_id	RegionOne
service_id	f71529314dab4a4d8eca427e701d209e
service_name	neutron
service_type	network
url	http://controller:9696

```
$ openstack endpoint create --region RegionOne \  
network admin http://controller:9696
```

Field	Value
enabled	True
id	1ee14289c9374dffb5db92a5c112fc4e
interface	admin
region	RegionOne
region_id	RegionOne
service_id	f71529314dab4a4d8eca427e701d209e
service_name	neutron
service_type	network
url	http://controller:9696

Configure networking options

You can deploy the Networking service using one of two architectures represented by options 1 and 2.

Option 1 deploys the simplest possible architecture that only supports attaching instances to provider (external) networks. No self-service (private) networks, routers, or floating IP addresses. Only the `admin` or other privileged user can manage provider networks.

- [Networking Option 1: Provider networks](#)

Here we choose Option 1.

Networking Option 1: Provider networks

Install and configure the Networking components on the *controller* node.

Install the components

```
# apt install neutron-server neutron-plugin-m12 \
neutron-linuxbridge-agent neutron-dhcp-agent \
neutron-metadata-agent -y
```

Configure the server component

The Networking server component configuration includes the database, authentication mechanism, message queue, topology change notifications, and plug-in.

- Edit the `/etc/neutron/neutron.conf` file and complete the following actions:
 - In the `[database]` section, configure database access:

```
[database]
...
connection = mysql+pymysql://neutron:NEUTRON_DBPASS@controller/neutron
```

Replace `NEUTRON_DBPASS` with the password you chose for the database.

Comment out or remove any other `connection` options in the `[database]` section.

- In the `[DEFAULT]` section, enable the Modular Layer 2 (ML2) plug-in and disable additional plug-ins:

```
[DEFAULT]
...
core_plugin = ml2
service_plugins =
```

- In the `[DEFAULT]` section, configure `RabbitMQ` message queue access:

```
[DEFAULT]
...
transport_url = rabbit://openstack:RABBIT_PASS@controller
```

Replace `RABBIT_PASS` with the password you chose for the `openstack` account in RabbitMQ.

- In the `[DEFAULT]` and `[keystone_authtoken]` sections, configure Identity service access:

```
[DEFAULT]
...
auth_strategy = keystone

[keystone_authtoken]
...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = Default
user_domain_name = Default
project_name = service
username = neutron
password = neutron
```

Replace `password = neutron` with the password you chose for the `neutron` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

- In the `[DEFAULT]` and `[nova]` sections, configure Networking to notify Compute of network topology changes:

```
[DEFAULT]
...
notify_nova_on_port_status_changes = True
notify_nova_on_port_data_changes = True

[nova]
...
auth_url = http://controller:35357
auth_type = password
project_domain_name = Default
user_domain_name = Default
region_name = RegionOne
project_name = service
username = nova
password = nova
```

Replace `password = nova` with the password you chose for the `nova` user in the Identity service.

Configure the Modular Layer 2 (ML2) plug-in

The ML2 plug-in uses the Linux bridge mechanism to build layer-2 (bridging and switching) virtual networking infrastructure for instances.

- Edit the `/etc/neutron/plugins/ml2/ml2_conf.ini` file and complete the following actions:
 - In the `[ml2]` section, enable flat and VLAN networks:

```
[m12]
...
type_drivers = flat,vlan
```

- In the `[m12]` section, disable self-service networks:

```
[m12]
...
tenant_network_types =
```

- In the `[m12]` section, enable the Linux bridge mechanism:

```
[m12]
...
mechanism_drivers = linuxbridge
```

After you configure the ML2 plug-in, removing values in the `type_drivers` option can lead to database inconsistency.

- In the `[m12]` section, enable the port security extension driver:

```
[m12]
...
extension_drivers = port_security
```

- In the `[m12_type_flat]` section, configure the provider virtual network as a flat network:

```
[m12_type_flat]
...
flat_networks = provider
```

- In the `[securitygroup]` section, enable [ipset](#) to increase efficiency of security group rules:

```
[securitygroup]
...
enable_ipset = True
```

Configure the Linux bridge agent

The Linux bridge agent builds layer-2 (bridging and switching) virtual networking infrastructure for instances and handles security groups.

- Edit the `/etc/neutron/plugins/ml2/linuxbridge_agent.ini` file and complete the following actions:
 - In the `[linux_bridge]` section, map the provider virtual network to the provider physical network interface:

```
[linux_bridge]
physical_interface_mappings = provider:PROVIDER_INTERFACE_NAME
```

Replace `PROVIDER_INTERFACE_NAME` with the name of the underlying provider physical network interface. See [Host networking](#) for more information.

in our case it is: `enp0s10`, the bridged nic of controller network.

- o In the `[vxlan]` section, disable VXLAN overlay networks:

```
[vxlan]
enable_vxlan = False
```

- o In the `[securitygroup]` section, enable security groups and configure the Linux bridge [iptables](#) firewall driver:

```
[securitygroup]
...
enable_security_group = True
firewall_driver = neutron.agent.linux.iptables_firewall.IptablesFirewallDriver
```

Configure the DHCP agent

The [DHCP agent](#) provides DHCP services for virtual networks.

- Edit the `/etc/neutron/dhcp_agent.ini` file and complete the following actions:
 - o In the `[DEFAULT]` section, configure the Linux bridge interface driver, Dnsmasq DHCP driver, and enable isolated metadata so instances on provider networks can access metadata over the network:

```
[DEFAULT]
...
interface_driver = neutron.agent.linux.interface.BridgeInterfaceDriver
dhcp_driver = neutron.agent.linux.dhcp.Dnsmasq
enable_isolated_metadata = True
```

Return to [Networking controller node configuration](#).

Configure the metadata agent

The [metadata agent](#) provides configuration information such as credentials to instances.

- Edit the `/etc/neutron/metadata_agent.ini` file and complete the following actions:
 - o In the `[DEFAULT]` section, configure the metadata host and shared secret:

```
[DEFAULT]
...
nova_metadata_ip = controller
metadata_proxy_shared_secret = METADATA_SECRET
```

Replace `METADATA_SECRET` with a suitable secret for the metadata proxy.

Configure the Compute service to use the Networking service

- Edit the `/etc/nova/nova.conf` file and perform the following actions:
 - In the `[neutron]` section, configure access parameters, enable the metadata proxy, and configure the secret:

```
[neutron]
...
url = http://controller:9696
auth_url = http://controller:35357
auth_type = password
project_domain_name = Default
user_domain_name = Default
region_name = RegionOne
project_name = service
username = neutron
password = neutron
service_metadata_proxy = True
metadata_proxy_shared_secret = METADATA_SECRET
```

Replace `password = neutron` with the password you chose for the `neutron` user in the Identity service.

Replace `METADATA_SECRET` with the secret you chose for the metadata proxy.

Finalize installation

1. Populate the database:

```
# su -s /bin/sh -c "neutron-db-manage --config-file /etc/neutron/neutron.conf \
--config-file /etc/neutron/plugins/ml2/ml2_conf.ini upgrade head" neutron
```

Database population occurs later for Networking because the script requires complete server and plug-in configuration files.

2. Restart the Compute API service:

```
# service nova-api restart
```

3. Restart the Networking services.

For both networking options:

```
# service neutron-server restart
# service neutron-linuxbridge-agent restart
# service neutron-dhcp-agent restart
# service neutron-metadata-agent restart
```

For networking option 2, also restart the layer-3 service:

```
# service neutron-l3-agent restart
```

Verify operation

```
root@controller:~# openstack network agent list --max-width 50
```

ID	Agent Type	Host	Availability Zone	Alive	State	Binary
1d661145-0941-411d-9b18-b3371fe57c4b	Linux bridge agent	control1er	None	True	UP	neutron-linuxbridge-agent
7502e1a3-998d-4aca-91e4-ca17e1b10c82	DHCP agent	control1er	nova	True	UP	neutron-dhcp-agent
7c47ac70-5de2-4442-8fc1-91fe97ae120f	Metadata agent	control1er	None	True	UP	neutron-metadata-agent

Neutron Install and configure compute node

The compute node handles connectivity and [security groups](#) for instances.

Install the components

```
# apt install neutron-linuxbridge-agent -y
```

Configure the common component

The Networking common component configuration includes the authentication mechanism, message queue, and plug-in.

- Edit the `/etc/neutron/neutron.conf` file and complete the following actions:
 - In the `[database]` section, comment out any `connection` options because compute nodes do not directly access the database.
 - In the `[DEFAULT]` section, configure `RabbitMQ` message queue access:

```
[DEFAULT]
...
transport_url = rabbit://openstack:RABBIT_PASS@controller
```

Replace `RABBIT_PASS` with the password you chose for the `openstack` account in RabbitMQ.

- In the `[DEFAULT]` and `[keystone_authtoken]` sections, configure Identity service access:

```
[DEFAULT]
...
auth_strategy = keystone

[keystone_authtoken]
...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = Default
user_domain_name = Default
project_name = service
username = neutron
password = neutron
```

Replace `password = neutron` with the password you chose for the `neutron` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

Configure networking options

Choose the same networking option that you chose for the controller node to configure services specific to it. Afterwards, return here and proceed to [Configure the Compute service to use the Networking service](#).

- [Networking Option 1: Provider networks](#)

Configure the Linux bridge agent

The Linux bridge agent builds layer-2 (bridging and switching) virtual networking infrastructure for instances and handles security groups.

- Edit the `/etc/neutron/plugins/ml2/linuxbridge_agent.ini` file and complete the following actions:
 - In the `[linux_bridge]` section, map the provider virtual network to the provider physical network interface:

```
[linux_bridge]
physical_interface_mappings = provider:PROVIDER_INTERFACE_NAME
```

Replace `PROVIDER_INTERFACE_NAME` with the name of the underlying provider physical network interface. See [Host networking](#) for more information.

- In the `[vxlan]` section, disable VXLAN overlay networks:

```
[vxlan]
enable_vxlan = False
```

- In the `[securitygroup]` section, enable security groups and configure the Linux bridge [iptables](#) firewall driver:

```
[securitygroup]
...
enable_security_group = True
firewall_driver = neutron.agent.linux.iptables_firewall.IptablesFirewallDriver
```

Return to [Networking compute node configuration](#).

Configure the Compute service to use the Networking service

- Edit the `/etc/nova/nova.conf` file and complete the following actions:
 - In the `[neutron]` section, configure access parameters:

```
[neutron]
...
url = http://controller:9696
auth_url = http://controller:35357
auth_type = password
project_domain_name = Default
user_domain_name = Default
region_name = RegionOne
project_name = service
username = neutron
password = neutron
```

Replace `password = neutron` with the password you chose for the `neutron` user in the Identity service.

Finalize installation

1. Restart the Compute service:

```
# service nova-compute restart
```

2. Restart the Linux bridge agent:

```
# service neutron-linuxbridge-agent restart
```

Verify operation

Perform these commands on the controller node.

1. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

2. List loaded extensions to verify successful launch of the `neutron-server` process:

```
$ neutron ext-list
```

alias	name
default-subnetpools	Default Subnetpools
network-ip-availability	Network IP Availability
network_availability_zone	Network Availability Zone
auto-allocated-topology	Auto Allocated Topology Services
ext-gw-mode	Neutron L3 Configurable external gateway mode
binding	Port Binding
agent	agent
subnet_allocation	Subnet Allocation
l3_agent_scheduler	L3 Agent Scheduler
tag	Tag support
external-net	Neutron external network
net-mtu	Network MTU
availability_zone	Availability Zone
quotas	Quota management support
l3-ha	HA Router extension
flavors	Neutron Service Flavors
provider	Provider Network
multi-provider	Multi Provider Network
address-scope	Address scope
extraroute	Neutron Extra Route
timestamp_core	Time Stamp Fields addition for core resources
router	Neutron L3 Router
extra_dhcp_opt	Neutron Extra DHCP opts
dns-integration	DNS Integration
security-group	security-group
dhcp_agent_scheduler	DHCP Agent Scheduler
router_availability_zone	Router Availability Zone
rbac-policies	RBAC Policies
standard-attr-description	standard-attr-description
port-security	Port Security
allowed-address-pairs	Allowed Address Pairs
dvr	Distributed Virtual Router

3. List agents to verify successful launch of the neutron agents:

```
$ openstack network agent list
```

```
root@controller:~# openstack network agent list --max-width 70
```

ID	Agent Type	Host	Availability Zone	Alive	State	Binary
143d7731-9227-4baf-9052-292d7aea6992	Linux bridge agent	compute	None	True	UP	neutron-libnxbridge-agent
1d661145-0941-411d-9b18-b3371fe57c4b	Linux bridge agent	control1er	None	True	UP	neutron-libnxbridge-agent
7502e1a3-998d-4aca-91e4-ca17e1b10c82	DHCP agent	control1er	nova	True	UP	neutron-dhcp-agent
7c47ac70-5de2-4442-8fc1-91fe97ae120f	Metadata agent	control1er	None	True	UP	neutron-metadata-agent

The output should indicate three agents on the controller node and one agent on each compute node.

Congratulations! Let's try booting an instance

Create provider network/subnetwork

ref: <https://docs.openstack.org/newton/install-guide-ubuntu/launch-instance-networks-provider.html>

The `--provider:physical_network` provider and `--provider:network_type flat` options connect the flat virtual network to the flat (native/untagged) physical network on the eth1 interface on the host

标注: 下边的创建网络里, 参数:

```
--provider-network-type flat \  
--provider-physical-network provider
```

对应的是:

```
/etc/neutron/plugins/ml2/ml2_conf.ini
```

```
[ml2_type_flat]
```

```
flat_networks = provider
```

```
/etc/neutron/plugins/ml2/linuxbridge_agent.ini
```

```
[linux_bridge]
```

```
physical_interface_mappings = provider:enp0s10
```

```

root@controller:~# . admin-openrc
root@controller:~# openstack network create --share --external \
> --provider-physical-network provider \
> --provider-network-type flat provider

```

Field	Value
admin_state_up	UP
availability_zone_hints	
availability_zones	
created_at	2017-08-23T17:14:21Z
description	
dns_domain	None
id	2a33434f-ba29-4645-9b5d-24f1509066f1
ipv4_address_scope	None
ipv6_address_scope	None
is_default	None
mtu	1500
name	provider
port_security_enabled	True
project_id	78c9c849237649a3a8c4526167427589
provider:network_type	flat
provider:physical_network	provider
provider:segmentation_id	None
qos_policy_id	None
revision_number	4
router:external	External
segments	None
shared	True
status	ACTIVE
subnets	
updated_at	2017-08-23T17:14:21Z

```

root@controller:~# neutron net-list

```

neutron CLI is deprecated and will be removed in the future. Use openstack CLI instead.

id	name	tenant_id	subnets
2a33434f-ba29-4645-9b5d-24f1509066f1	provider	78c9c849237649a3a8c4526167427589	

```

root@controller:~# openstack network list

```

ID	Name	Subnets
2a33434f-ba29-4645-9b5d-24f1509066f1	provider	

```

root@controller:~# openstack subnet create --network provider \
> --allocation-pool start=146.11.41.230,end=146.11.41.233 \
> --dns-nameserver 147.128.5.12 --gateway 146.11.40.1 \
> --subnet-range 146.11.40.1/23 provider

```

Field	Value
-------	-------


```

+-----+
| allocation_pools | 146.11.41.230-146.11.41.233 |
| cidr             | 146.11.40.0/23             |
| created_at       | 2017-08-23T17:17:54Z       |
| description      |                             |
| dns_nameservers  | 147.128.5.12               |
| enable_dhcp      | True                        |
| gateway_ip       | 146.11.40.1                 |
| host_routes      |                             |
| id               | 9b118521-59b5-40ee-a439-9d59c3b392ea |
| ip_version       | 4                           |
| ipv6_address_mode | None                        |
| ipv6_ra_mode     | None                        |
| name             | provider                    |
| network_id       | 2a33434f-ba29-4645-9b5d-24f1509066f1 |
| project_id       | 78c9c849237649a3a8c4526167427589 |
| revision_number  | 2                            |
| segment_id       | None                        |
| service_types    |                             |
| subnetpool_id    | None                        |
| updated_at       | 2017-08-23T17:17:54Z       |
+-----+

```

Create flavor

The smallest default flavor consumes 512 MB memory per instance. For environments with compute nodes containing less than 4 GB memory, we recommend creating the `m1.nano` flavor that only requires 64 MB per instance. Only use this flavor with the CirrOS image for testing purposes.

```
$ openstack flavor create --id 0 --vcpus 1 --ram 64 --disk 1 m1.nano
```

Field	Value
OS-FLV-DISABLED:disabled	False
OS-FLV-EXT-DATA:ephemeral	0
disk	1
id	0
name	m1.nano
os-flavor-access:is_public	True
ram	64
rxtx_factor	1.0
swap	
vcpus	1

Add security group rules

By default, the `default` security group applies to all instances and includes firewall rules that deny remote access to instances. For Linux images such as CirrOS, we recommend allowing at least ICMP (ping) and secure shell (SSH).

- Add rules to the `default` security group:
 - Permit [ICMP](#) (ping):

```
$ openstack security group rule create --proto icmp default
```

Field	Value
created_at	2016-10-05T09:52:31Z
description	
direction	ingress
ethertype	IPv4
headers	
id	6ee8d630-9803-4d3d-9aea-8c795abbedc2
port_range_max	None
port_range_min	None
project_id	77ae8d7104024123af342ffb0a6f1d88
project_id	77ae8d7104024123af342ffb0a6f1d88
protocol	icmp
remote_group_id	None
remote_ip_prefix	0.0.0.0/0
revision_number	1
security_group_id	4ceee3d4-d2fe-46c1-895c-382033e87b0d
updated_at	2016-10-05T09:52:31Z

- o Permit secure shell (SSH) access:

```
$ openstack security group rule create --proto tcp --dst-port 22 default
```

Field	Value
created_at	2016-10-05T09:54:50Z
description	
direction	ingress
ethertype	IPv4
headers	
id	3cd0a406-43df-4741-ab29-b5e7dcb7469d
port_range_max	22
port_range_min	22
project_id	77ae8d7104024123af342ffb0a6f1d88
project_id	77ae8d7104024123af342ffb0a6f1d88
protocol	tcp
remote_group_id	None
remote_ip_prefix	0.0.0.0/0
revision_number	1
security_group_id	4ceee3d4-d2fe-46c1-895c-382033e87b0d
updated_at	2016-10-05T09:54:50Z

Launch an instance

ref: [Launch an instance on the provider network](#)

Determine instance options

To launch an instance, you must at least specify the flavor, image name, network, security group, key, and instance name.

1. On the controller node, source the `demo` credentials to gain access to user-only CLI commands:

```
$ . demo-openrc
```

2. A flavor specifies a virtual resource allocation profile which includes processor, memory, and storage.
List available flavors:

```
$ openstack flavor list

+-----+-----+-----+-----+-----+-----+-----+
| ID | Name   | RAM | Disk | Ephemeral | VCPUs | Is Public |
+-----+-----+-----+-----+-----+-----+-----+
| 0  | m1.nano | 64  | 1    | 0          | 1     | True      |
+-----+-----+-----+-----+-----+-----+-----+
```

You can also reference a flavor by ID.

3. List available images:

```
$ openstack image list

+-----+-----+-----+
| ID | Name | Status |
+-----+-----+-----+
| 390eb5f7-8d49-41ec-95b7-68c0d5d54b34 | cirros | active |
+-----+-----+-----+
```

This instance uses the `cirros` image.

4. List available networks:

```
$ openstack network list

+-----+-----+-----+
-+
| ID | Name | Subnets |
|
+-----+-----+-----+
-+
| 4716ddfe-6e60-40e7-b2a8-42e57bf3c31c | selfservice | 2112d5eb-f9d6-45fd-906e-7cabd38b7c7c |
|
| b5b6993c-ddf9-40e7-91d0-86806a42edb8 | provider    | 310911f6-acf0-4a47-824e-3032916582ff |
|
+-----+-----+-----+
-+
```

This instance uses the `provider` provider network. However, you must reference this network using the ID instead of the name.

5. List available security groups:

```
$ openstack security group list

+-----+-----+-----+-----+
| ID                                     | Name   | Description          | Project |
+-----+-----+-----+-----+
| dd2b614c-3dad-48ed-958b-b155a3b38515 | default | Default security group |         |
a516b957032844328896baa01e0f906c |         |                         |         |
+-----+-----+-----+-----+
```

This instance uses the `default` security group.

Launch the instance

1. Launch the instance:

Replace `PROVIDER_NET_ID` with the ID of the `provider` provider network.

If you chose option 1 and your environment contains only one network, you can omit the `--nic` option because OpenStack automatically chooses the only network available.

```
root@controller:~# openstack server create --flavor m1.nano --image cirros \
> --nic net-id=2a33434f-ba29-4645-9b5d-24f1509066f1 --security-group default provider-
instance
```

Field	Value
OS-DCF:diskConfig	MANUAL
OS-EXT-AZ:availability_zone	
OS-EXT-STS:power_state	NOSTATE
OS-EXT-STS:task_state	scheduling
OS-EXT-STS:vm_state	building
OS-SRV-USG:launched_at	None
OS-SRV-USG:terminated_at	None
accessIPv4	
accessIPv6	
addresses	
adminPass	MnjXdXf3qHia
config_drive	
created	2017-08-23T17:29:04Z
flavor	m1.nano (0)
hostId	
id	02f54ef9-e867-4c1a-88f9-8eddd144da6f
image	cirros (c17e391e-93e1-4480-9cf3-bf8623063e61)
key_name	None
name	provider-instance
progress	0
project_id	cb015df53fb34d90b077e4c36ce35826
properties	
security_groups	name='default'
status	BUILD
updated	2017-08-23T17:29:05Z
user_id	cb98fad69e84459bb48f42130d5c0ce5
volumes_attached	

2. Check the status of your instance:

```
root@controller:~# nova list
```

```
+-----+-----+-----+-----+-----+
+-----+
| ID                      | Name              | Status | Task State | Power
State | Networks |
+-----+-----+-----+-----+-----+
+-----+
| 02f54ef9-e867-4c1a-88f9-8eddd144da6f | provider-instance | BUILD  | scheduling | NOSTATE
|          |
+-----+-----+-----+-----+-----+
+-----+
```

```
root@controller:~# openstack server list
```

```
+-----+-----+-----+-----+-----+
+
| ID                      | Name              | Status | Networks | Image Name
|
+-----+-----+-----+-----+-----+
+
| 02f54ef9-e867-4c1a-88f9-8eddd144da6f | provider-instance | BUILD  |          | cirros
|
+-----+-----+-----+-----+-----+
+
```

The status changes from `BUILD` to `ACTIVE` when the build process successfully completes.

Access the instance using the virtual console

1. Obtain a [Virtual Network Computing \(VNC\)](#) session URL for your instance and access it from a web browser:

```
$ openstack console url show provider-instance
```

```
+-----+-----+-----+-----+-----+
| Field | Value
+-----+-----+-----+-----+-----+
| type  | novnc
| url   | http://controller:6080/vnc_auto.html?token=5eeccb47-525c-4918-ac2a-3ad1e9f1f493
+-----+-----+-----+-----+-----+
```

If your web browser runs on a host that cannot resolve the `controller` host name, you can replace `controller` with the IP address of the management interface on your controller node.

The CirrOS image includes conventional user name/password authentication and provides these credentials at the login prompt. After logging into CirrOS, we recommend that you verify network connectivity using `ping`.

2. Verify access to the provider physical network gateway:

```
$ ping -c 4 203.0.113.1

PING 203.0.113.1 (203.0.113.1) 56(84) bytes of data.
64 bytes from 203.0.113.1: icmp_req=1 ttl=64 time=0.357 ms
64 bytes from 203.0.113.1: icmp_req=2 ttl=64 time=0.473 ms
64 bytes from 203.0.113.1: icmp_req=3 ttl=64 time=0.504 ms
64 bytes from 203.0.113.1: icmp_req=4 ttl=64 time=0.470 ms

--- 203.0.113.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 2998ms
rtt min/avg/max/mdev = 0.357/0.451/0.504/0.055 ms
```

3. Verify access to the internet:

```
$ ping -c 4 openstack.org

PING openstack.org (174.143.194.225) 56(84) bytes of data.
64 bytes from 174.143.194.225: icmp_req=1 ttl=53 time=17.4 ms
64 bytes from 174.143.194.225: icmp_req=2 ttl=53 time=17.5 ms
64 bytes from 174.143.194.225: icmp_req=3 ttl=53 time=17.7 ms
64 bytes from 174.143.194.225: icmp_req=4 ttl=53 time=17.5 ms

--- openstack.org ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3003ms
rtt min/avg/max/mdev = 17.431/17.575/17.734/0.143 ms
```

Access the instance remotely

1. Verify connectivity to the instance from the controller node or any host on the provider physical network:

```
$ ping -c 4 203.0.113.103

PING 203.0.113.103 (203.0.113.103) 56(84) bytes of data.
64 bytes from 203.0.113.103: icmp_req=1 ttl=63 time=3.18 ms
64 bytes from 203.0.113.103: icmp_req=2 ttl=63 time=0.981 ms
64 bytes from 203.0.113.103: icmp_req=3 ttl=63 time=1.06 ms
64 bytes from 203.0.113.103: icmp_req=4 ttl=63 time=0.929 ms

--- 203.0.113.103 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3002ms
rtt min/avg/max/mdev = 0.929/1.539/3.183/0.951 ms
```

2. Access your instance using SSH from the controller node or any host on the provider physical network:


```
$ ssh cirros@203.0.113.103
```

```
The authenticity of host '203.0.113.102 (203.0.113.102)' can't be established.  
RSA key fingerprint is ed:05:e9:e7:52:a0:ff:83:68:94:c7:d1:f2:f8:e2:e9.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added '203.0.113.102' (RSA) to the list of known hosts.
```

If your instance does not launch or seem to work as you expect, see the [Instance Boot Failures](#) section in OpenStack Operations Guide for more information or use one of the [many other options](#) to seek assistance. We want your first installation to work!

Return to [Launch an instance](#).

[ISSUE] DHCP failure in VM troubleshooting

ref: <https://docs.openstack.org/neutron/pike/admin/intro-basic-networking.html>

what it is like

in VM console (initial dhcp discover)

```
$ ifup eth0  
udhcpd (v1.20.1) started  
Sending discover...  
Sending discover...  
Sending discover...  
Usage: /sbin/cirros-dhcpd <up|down>  
No lease, failing
```

in controller console (monitor log)

```
root@controller:~# tail -f /var/log/syslog  
Aug 24 03:06:30 controller dhclient[1166]: DHCPREQUEST of 146.11.41.129 on enp0s10 to  
147.128.5.12 port 67 (xid=0x5d38ef7e)  
Aug 24 03:06:30 controller dhclient[1166]: DHCPACK of 146.11.41.129 from 147.128.5.12  
Aug 24 03:06:30 controller dhclient[1166]: Invalid domain list.  
Aug 24 03:06:30 controller dhclient[1166]: suspect value in domain_search option - discarded  
Aug 24 03:06:30 controller dhclient[1166]: Invalid domain list.  
Aug 24 03:06:30 controller dhclient[1166]: suspect value in domain_search option - discarded  
Aug 24 03:06:30 controller dhclient[1166]: Invalid domain list.  
Aug 24 03:06:30 controller dhclient[1166]: bound to 146.11.41.129 -- renewal in 12824 seconds.  
Aug 24 03:12:18 controller dnsmasq-dhcp[18894]: DHCPDISCOVER(ns-a6e0220e-ec) fa:16:3e:bb:c6:13  
Aug 24 03:12:18 controller dnsmasq-dhcp[18894]: DHCPOFFER(ns-a6e0220e-ec) 146.11.41.232  
fa:16:3e:bb:c6:13  
Aug 24 03:13:19 controller dnsmasq-dhcp[18894]: DHCPDISCOVER(ns-a6e0220e-ec) fa:16:3e:bb:c6:13  
Aug 24 03:13:19 controller dnsmasq-dhcp[18894]: DHCPOFFER(ns-a6e0220e-ec) 146.11.41.232  
fa:16:3e:bb:c6:13  
Aug 24 03:14:19 controller dnsmasq-dhcp[18894]: DHCPDISCOVER(ns-a6e0220e-ec) fa:16:3e:bb:c6:13  
Aug 24 03:14:19 controller dnsmasq-dhcp[18894]: DHCPOFFER(ns-a6e0220e-ec) 146.11.41.232  
fa:16:3e:bb:c6:13
```

By tcpdump from controller bridge, it's found the DHCPOFFER was sent to VM:

```
root@controller:~# tcpdump -i brq2a33434f-ba -vv port 67 or port 68 -e -n
tcpdump: listening on brq2a33434f-ba, link-type EN10MB (Ethernet), capture size 262144 bytes
04:04:22.830138 fa:16:3e:bb:c6:13 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800), length 332: (tos
0x0, ttl 64, id 0, offset 0, flags [none], proto UDP (17), length 318)
    0.0.0.0.68 > 255.255.255.255.67: [udp sum ok] BOOTP/DHCP, Request from fa:16:3e:bb:c6:13,
length 290, xid 0x1fac2751, Flags [none] (0x0000)
        Client-Ethernet-Address fa:16:3e:bb:c6:13
        Vendor-rfc1048 Extensions
            Magic Cookie 0x63825363
            DHCP-Message Option 53, length 1: Discover
            Client-ID Option 61, length 7: ether fa:16:3e:bb:c6:13
            MSZ Option 57, length 2: 576
            Parameter-Request Option 55, length 9:
                Subnet-Mask, Default-Gateway, Domain-Name-Server, Hostname
                Domain-Name, MTU, BR, NTP
                Classless-Static-Route
            Vendor-Class Option 60, length 12: "udhcp 1.20.1"
            Hostname Option 12, length 6: "cirros"
04:04:22.831801 fa:16:3e:8b:53:5e > fa:16:3e:bb:c6:13, ethertype IPv4 (0x0800), length 370: (tos
0xc0, ttl 64, id 3044, offset 0, flags [none], proto UDP (17), length 356)
    146.11.41.230.67 > 146.11.41.232.68: [udp sum ok] BOOTP/DHCP, Reply, length 328, xid
0x1fac2751, Flags [none] (0x0000)
        Your-IP 146.11.41.232
        Server-IP 146.11.41.230
        Client-Ethernet-Address fa:16:3e:bb:c6:13
        Vendor-rfc1048 Extensions
            Magic Cookie 0x63825363
            DHCP-Message Option 53, length 1: Offer
            Server-ID Option 54, length 4: 146.11.41.230
            Lease-Time Option 51, length 4: 86400
            RN Option 58, length 4: 43200
            RB Option 59, length 4: 75600
            Subnet-Mask Option 1, length 4: 255.255.254.0
            BR Option 28, length 4: 146.11.41.255
            Domain-Name Option 15, length 14: "openstacklocal"
            Default-Gateway Option 3, length 4: 146.11.40.1
            Classless-Static-Route Option 121, length 14: (169.254.169.254/32:146.11.41.230),
(default:146.11.40.1)
            Domain-Name-Server Option 6, length 4: 147.128.5.12
            MTU Option 26, length 2: 1500
04:04:52.504181 08:2e:5f:5d:63:00 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800), length 358: (tos
0x0, ttl 120, id 5320, offset 0, flags [DF], proto UDP (17), length 344)
```

While from compute , tcpdump the br-int bridge shows it's not received

```

root@compute:~# tcpdump -i brq2a33434f-ba -vv port 67 or port 68 -e -n
tcpdump: listening on brq2a33434f-ba, link-type EN10MB (Ethernet), capture size 262144 bytes
03:51:04.456668 fa:16:3e:bb:c6:13 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800), length 332: (tos
0x0, ttl 64, id 0, offset 0, flags [none], proto UDP (17), length 318)
    0.0.0.0.68 > 255.255.255.255.67: [udp sum ok] BOOTP/DHCP, Request from fa:16:3e:bb:c6:13,
length 290, xid 0xe5e8f024, Flags [none] (0x0000)
        Client-Ethernet-Address fa:16:3e:bb:c6:13
        Vendor-rfc1048 Extensions
            Magic Cookie 0x63825363
            DHCP-Message Option 53, length 1: Discover
            Client-ID Option 61, length 7: ether fa:16:3e:bb:c6:13
            MSZ Option 57, length 2: 576
            Parameter-Request Option 55, length 9:
                Subnet-Mask, Default-Gateway, Domain-Name-Server, Hostname
                Domain-Name, MTU, BR, NTP
                Classless-Static-Route
            Vendor-Class Option 60, length 12: "udhcp 1.20.1"
            Hostname Option 12, length 6: "cirros"
03:51:30.022360 08:2e:5f:5d:63:00 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800), length 358: (tos
0x0, ttl 120, id 29901, offset 0, flags [DF], proto UDP (17), length 344)
    147.128.5.12.67 > 255.255.255.255.68: [udp sum ok] BOOTP/DHCP, Reply, length 316, xid
0x4a42e788, Flags [Broadcast] (0x8000)
        Client-IP 146.11.40.250
        Gateway-IP 146.11.40.1
        Client-Ethernet-Address d0:bf:9c:df:7a:a5
        Vendor-rfc1048 Extensions
            Magic Cookie 0x63825363
            DHCP-Message Option 53, length 1: ACK
            Server-ID Option 54, length 4: 147.128.5.12
            Subnet-Mask Option 1, length 4: 255.255.254.0
            Vendor-Option Option 43, length 5: 220.3.78.65.80
            Domain-Name Option 15, length 18: "cn.ao.ericsson.se^@"
            Default-Gateway Option 3, length 4: 146.11.40.1
            Domain-Name-Server Option 6, length 12: 147.128.5.12,193.181.14.11,193.181.14.10
            Netbios-Name-Server Option 44, length 8: 146.11.115.50,146.11.116.30
            Netbios-Node Option 46, length 1: h-node
03:51:30.023490 2c:76:8a:1f:47:00 > ff:ff:ff:ff:ff:ff, ethertype IPv4 (0x0800), length 358: (tos
0x0, ttl 119, id 2187, offset 0, flags [DF], proto UDP (17), length 344)

```

Conclusion:

the DHCP offer was sent out from DHCP agent dnsmasq, but the package cannot be captured from host bridge connecting to vm eth0. the issue is located in the provider network router, the ECN router 146.11.40.1 in our office.

By searching online, there is a tech called DHCP snooping to prevent multiple dhcp server in one LAN from router, which makes sense.

Cinder

Cinder on controller

Here we provide a iSCSI driver backend cinder practice

ref: <https://docs.openstack.org/ocata/install-guide-ubuntu/cinder.html>

The OpenStack Block Storage service (cinder) adds persistent storage to a virtual machine. Block Storage provides an infrastructure for managing volumes, and interacts with OpenStack Compute to provide volumes for instances. The service also enables management of volume snapshots, and volume types.

The Block Storage service consists of the following components:

- cinder-api
Accepts API requests, and routes them to the `cinder-volume` for action.
- cinder-volume
Interacts directly with the Block Storage service, and processes such as the `cinder-scheduler`. It also interacts with these processes through a message queue. The `cinder-volume` service responds to read and write requests sent to the Block Storage service to maintain state. It can interact with a variety of storage providers through a driver architecture.
- cinder-scheduler daemon
Selects the optimal storage provider node on which to create the volume. A similar component to the `nova-scheduler`.
- cinder-backup daemon
The `cinder-backup` service provides backing up volumes of any type to a backup storage provider. Like the `cinder-volume` service, it can interact with a variety of storage providers through a driver architecture.
- Messaging queue
Routes information between the Block Storage processes.

Install and configure controller node

This section describes how to install and configure the Block Storage service, code-named cinder, on the controller node. This service requires at least one additional storage node that provides volumes to instances.

Prerequisites

Before you install and configure the Block Storage service, you must create a database, service credentials, and API endpoints.

1. To create the database, complete these steps:

- Use the database access client to connect to the database server as the `root` user:

```
# mysql
```

- Create the `cinder` database:

```
MariaDB [(none)]> CREATE DATABASE cinder;
```

- Grant proper access to the `cinder` database:

```
MariaDB [(none)]> GRANT ALL PRIVILEGES ON cinder.* TO 'cinder'@'localhost' \
IDENTIFIED BY 'CINDER_DBPASS';
MariaDB [(none)]> GRANT ALL PRIVILEGES ON cinder.* TO 'cinder'@'%' \
IDENTIFIED BY 'CINDER_DBPASS';
```

Replace `CINDER_DBPASS` with a suitable password.

- Exit the database access client.

2. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

3. To create the service credentials, complete these steps:

- Create a `cinder` user:

```
$ openstack user create --domain default --password-prompt cinder
```

User Password:

Repeat User Password:

Field	Value
domain_id	default
enabled	True
id	9d7e33de3e1a498390353819bc7d245d
name	cinder
options	{}
password_expires_at	None

- Add the `admin` role to the `cinder` user:

```
$ openstack role add --project service --user cinder admin
```

This command provides no output.

- Create the `cinderv2` and `cinderv3` service entities:

```
$ openstack service create --name cinderv2 \
--description "OpenStack Block Storage" volumev2
```

Field	Value
description	OpenStack Block Storage
enabled	True
id	eb9fd245bdbc414695952e93f29fe3ac
name	cinderv2
type	volumev2

```
$ openstack service create --name cinderv3 \
  --description "OpenStack Block Storage" volumev3
```

```
+-----+-----+
| Field      | Value                                |
+-----+-----+
| description | OpenStack Block Storage             |
| enabled     | True                                |
| id          | ab3bbbef780845a1a283490d281e7fda  |
| name        | cinderv3                            |
| type        | volumev3                            |
+-----+-----+
```

The Block Storage services require two service entities.

4. Create the Block Storage service API endpoints:

```
$ openstack endpoint create --region RegionOne \  
  volumev2 public http://controller:8776/v2/%(project_id)s
```

Field	Value
enabled	True
id	513e73819e14460fb904163f41ef3759
interface	public
region	RegionOne
region_id	RegionOne
service_id	eb9fd245bdbc414695952e93f29fe3ac
service_name	cinderv2
service_type	volumev2
url	http://controller:8776/v2/%(project_id)s

```
$ openstack endpoint create --region RegionOne \  
  volumev2 internal http://controller:8776/v2/%(project_id)s
```

Field	Value
enabled	True
id	6436a8a23d014cfdb69c586eff146a32
interface	internal
region	RegionOne
region_id	RegionOne
service_id	eb9fd245bdbc414695952e93f29fe3ac
service_name	cinderv2
service_type	volumev2
url	http://controller:8776/v2/%(project_id)s

```
$ openstack endpoint create --region RegionOne \  
  volumev2 admin http://controller:8776/v2/%(project_id)s
```

Field	Value
enabled	True
id	e652cf84dd334f359ae9b045a2c91d96
interface	admin
region	RegionOne
region_id	RegionOne
service_id	eb9fd245bdbc414695952e93f29fe3ac
service_name	cinderv2
service_type	volumev2
url	http://controller:8776/v2/%(project_id)s

```
$ openstack endpoint create --region RegionOne \  
  volumev3 public http://controller:8776/v3/%(project_id)s
```

Field	Value
enabled	True
id	03fa2c90153546c295bf30ca86b1344b
interface	public
region	RegionOne
region_id	RegionOne
service_id	ab3bbbef780845a1a283490d281e7fda
service_name	cinderv3
service_type	volumev3
url	http://controller:8776/v3/%(project_id)s

```
$ openstack endpoint create --region RegionOne \  
  volumev3 internal http://controller:8776/v3/%(project_id)s
```

Field	Value
enabled	True
id	94f684395d1b41068c70e4ecb11364b2
interface	internal
region	RegionOne
region_id	RegionOne
service_id	ab3bbbef780845a1a283490d281e7fda
service_name	cinderv3
service_type	volumev3
url	http://controller:8776/v3/%(project_id)s

```
$ openstack endpoint create --region RegionOne \  
  volumev3 admin http://controller:8776/v3/%(project_id)s
```

Field	Value
enabled	True
id	4511c28a0f9840c78bacb25f10f62c98
interface	admin
region	RegionOne
region_id	RegionOne
service_id	ab3bbbef780845a1a283490d281e7fda
service_name	cinderv3
service_type	volumev3
url	http://controller:8776/v3/%(project_id)s

The Block Storage services require endpoints for each service entity.

Install and configure components

Install the packages:

```
# apt install cinder-api cinder-scheduler
```

Edit the `/etc/cinder/cinder.conf` file and complete the following actions:

- In the `[database]` section, configure database access:

```
[database]
# ...
connection = mysql+pymysql://cinder:CINDER_DBPASS@controller/cinder
```

Replace `CINDER_DBPASS` with the password you chose for the Block Storage database.

- In the `[DEFAULT]` section, configure `RabbitMQ` message queue access:

```
[DEFAULT]
# ...
transport_url = rabbit://openstack:RABBIT_PASS@controller
```

Replace `RABBIT_PASS` with the password you chose for the `openstack` account in `RabbitMQ`.

- In the `[DEFAULT]` and `[keystone_authtoken]` sections, configure Identity service access:

```
[DEFAULT]
# ...
auth_strategy = keystone

[keystone_authtoken]
# ...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = default
user_domain_name = default
project_name = service
username = cinder
password = cinder
```

Replace `password` with the password you chose for the `cinder` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

- In the `[DEFAULT]` section, configure the `my_ip` option to use the management interface IP address of the controller node:

```
[DEFAULT]
# ...
my_ip = 10.20.0.10
```

- In the `[oslo_concurrency]` section, configure the lock path:

```
[oslo_concurrency]
# ...
lock_path = /var/lib/cinder/tmp
```

Populate the Block Storage database:

```
# su -s /bin/sh -c "cinder-manage db sync" cinder
```

Ignore any deprecation messages in this output.

Configure Compute to use Block Storage

- Edit the `/etc/nova/nova.conf` file and add the following to it:

```
[cinder]
os_region_name = RegionOne
```

Finalize installation

1. Restart the Compute API service:

```
# service nova-api restart
```

2. Restart the Block Storage services:

```
# service cinder-scheduler restart
# service apache2 restart
```

Cinder on block storage node

configure storage network for compute

Check nic name

```
root@compute:~# dmesg | grep renamed
[ 2.730898] e1000 0000:00:09.0 enp0s9: renamed from eth2
[ 2.731826] e1000 0000:00:08.0 enp0s8: renamed from eth1
[ 2.732819] e1000 0000:00:0a.0 enp0s10: renamed from eth3
[ 2.735645] e1000 0000:00:03.0 enp0s3: renamed from eth0
```

`eth2` was named as `enp0s9`

Edit `/etc/network/interfaces`

```
# storage network eth2
auto enp0s9
iface enp0s9 inet static
address 192.168.199.20
netmask 255.255.255.0
```

```
# ifup enp0s9
```

Create cinder machine: storage

Storage actions & cinder on storage node

Clone it from base VM and add a virtual disk for storage vm

Management net eth0 (enp0s3) and storage net eth2 (enp0s9)

Edit `/etc/network/interfaces`

```
# management network eth0

auto enp0s3
iface enp0s3 inet static
address 10.20.0.30
netmask 255.255.255.0
```

```
# storage network eth2
auto enp0s9
iface enp0s9 inet static
address 192.168.199.30
netmask 255.255.255.0
```

```
//start the two nics
# ifup enp0s3
# ifup enp0s9
```

configure NTP by editing `/etc/chrony/chrony.conf`

```
server 10.20.0.10 iburst
```

change hostname and hosts

```
# echo storage > /etc/hostname
# echo 10.20.0.10    controller >> /etc/hosts
# echo 10.20.0.20    compute >> /etc/hosts
# hostname storage
```

Check new disk was there already

check by `fdisk -l` , /dev/sdb is there :-).

```
root@storage:~# fdisk -l
Disk /dev/sda: 50 GiB, 53687091200 bytes, 104857600 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x3ce50a75

Device      Boot  Start        End  Sectors  Size Id Type
/dev/sda1   *      2048      999423    997376   487M 83 Linux
/dev/sda2             1001470 104855551 103854082  49.5G  5 Extended
/dev/sda5             1001472 104855551 103854080  49.5G 8e Linux LVM

Disk /dev/sdb: 50 GiB, 53687091200 bytes, 104857600 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/mapper/ubuntu--vg-root: 45.5 GiB, 48876224512 bytes, 95461376 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes

Disk /dev/mapper/ubuntu--vg-swap_1: 4 GiB, 4294967296 bytes, 8388608 sectors
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
```

Install and configure a storage node

This section describes how to install and configure storage nodes for the Block Storage service. For simplicity, this configuration references one storage node with an empty local block storage device. The instructions use `/dev/sdb` , but you can substitute a different value for your particular node.

The service provisions logical volumes on this device using the [LVM](#) driver and provides them to instances via [iSCSI](#) transport. You can follow these instructions with minor modifications to horizontally scale your environment with additional storage nodes.

Prerequisites

Before you install and configure the Block Storage service on the storage node, you must prepare the storage device.

Perform these steps on the storage node.

1. Install the supporting utility packages:

```
# apt install lvm2
```

Some distributions include LVM by default.

2. Create the LVM physical volume `/dev/sdb`:

```
# pvcreate /dev/sdb

Physical volume "/dev/sdb" successfully created
```

3. Create the LVM volume group `cinder-volumes`:

```
# vgcreate cinder-volumes /dev/sdb

Volume group "cinder-volumes" successfully created
```

The Block Storage service creates logical volumes in this volume group.

4. Only instances can access Block Storage volumes. However, the underlying operating system manages the devices associated with the volumes. By default, the LVM volume scanning tool scans the `/dev` directory for block storage devices that contain volumes. If projects use LVM on their volumes, the scanning tool detects these volumes and attempts to cache them which can cause a variety of problems with both the underlying operating system and project volumes. You must reconfigure LVM to scan only the devices that contain the `cinder-volumes` volume group. Edit the `/etc/lvm/lvm.conf` file and complete the following actions:

- In the `devices` section, add a filter that accepts the `/dev/sdb` device and rejects all other devices:

```
devices {
...
filter = [ "a/sdb/", "r./.*/" ]
```

Each item in the filter array begins with `a` for **accept** or `r` for **reject** and includes a regular expression for the device name. The array must end with `r./.*/` to reject any remaining devices. You can use the `vgscan --vvvv` command to test filters.

If your storage nodes use LVM on the operating system disk, you must also add the associated device to the filter. For example, if the `/dev/sda` device contains the operating system:

```
filter = [ "a/sda/", "a/sdb/", "r./.*/" ]
```

Similarly, if your compute nodes use LVM on the operating system disk, you must also modify the filter in the `/etc/lvm/lvm.conf` file on those nodes to include only the operating system disk. For example, if the `/dev/sda` device contains the operating system:

```
filter = [ "a/sda/", "r./.*/" ]
```

Install and configure components

Install the packages:

```
# apt install cinder-volume -y
```

1. Edit the `/etc/cinder/cinder.conf` file and complete the following actions:

- o In the `[database]` section, configure database access:

```
[database]
# ...
connection = mysql+pymysql://cinder:CINDER_DBPASS@controller/cinder
```

Replace `CINDER_DBPASS` with the password you chose for the Block Storage database.

- o In the `[DEFAULT]` section, configure `RabbitMQ` message queue access:

```
[DEFAULT]
# ...
transport_url = rabbit://openstack:RABBIT_PASS@controller
```

Replace `RABBIT_PASS` with the password you chose for the `openstack` account in `RabbitMQ`.

- o In the `[DEFAULT]` and `[keystone_authtoken]` sections, configure Identity service access:

```
[DEFAULT]
# ...
auth_strategy = keystone

[keystone_authtoken]
# ...
auth_uri = http://controller:5000
auth_url = http://controller:35357
memcached_servers = controller:11211
auth_type = password
project_domain_name = default
user_domain_name = default
project_name = service
username = cinder
password = cinder
```

Replace `password` with the password you chose for the `cinder` user in the Identity service.

Comment out or remove any other options in the `[keystone_authtoken]` section.

- o In the `[DEFAULT]` section, configure the `my_ip` option:

```
[DEFAULT]
# ...
my_ip = STORAGE_INTERFACE_IP_ADDRESS
```

Replace `STORAGE_INTERFACE_IP_ADDRESS` with the IP address of the storage network (eth2).

- o In the `[lvm]` section, configure the LVM back end with the LVM driver, `cinder-volumes` volume group, iSCSI protocol, and appropriate iSCSI service:

```
[lvm]
# ...
volume_driver = cinder.volume.drivers.lvm.LVMVolumeDriver
volume_group = cinder-volumes
iscsi_protocol = iscsi
iscsi_helper = tgtadm
```

- In the `[DEFAULT]` section, enable the LVM back end:

```
[DEFAULT]
# ...
enabled_backends = lvm
```

Back-end names are arbitrary. As an example, this guide uses the name of the driver as the name of the back end.

- In the `[DEFAULT]` section, configure the location of the Image service API:

```
[DEFAULT]
# ...
glance_api_servers = http://controller:9292
```

- In the `[oslo_concurrency]` section, configure the lock path:

```
[oslo_concurrency]
# ...
lock_path = /var/lib/cinder/tmp
```

Finalize installation

1. Restart the Block Storage volume service including its dependencies:

```
# service tgt restart
# service cinder-volume restart
```

Verify operation

Verify operation of the Block Storage service.

Perform these commands on the controller node.

1. Source the `admin` credentials to gain access to admin-only CLI commands:

```
$ . admin-openrc
```

2. List service components to verify successful launch of each process:

```
root@controller:~# openstack volume service list
```

Binary	Host	Zone	Status	State	Updated At
cinder-scheduler	controller	nova	enabled	up	2017-08-24T14:35:44.000000
cinder-volume	storage@lvm	nova	enabled	up	2017-08-24T14:35:39.000000

Let's try something on block storage!

Create a volume

1. Source the `demo` credentials to perform the following steps as a non-administrative project:

```
$ . demo-openrc
```

2. Create a 1 GB volume:

```
$ openstack volume create --size 1 volume1
```

Field	Value
attachments	[]
availability_zone	nova
bootable	false
consistencygroup_id	None
created_at	2016-03-08T14:30:48.391027
description	None
encrypted	False
id	a1e8be72-a395-4a6f-8e07-856a57c39524
multiattach	False
name	volume1
properties	
replication_status	disabled
size	1
snapshot_id	None
source_volid	None
status	creating
type	None
updated_at	None
user_id	684286a9079845359882afc3aa5011fb

3. After a short time, the volume status should change from `creating` to `available`:


```
root@controller:~# openstack volume list
```

ID	Display Name	Status	Size	Attached to
81ffed40-ed71-495d-bfa9-8fb8c72cf222	volume1	available	1	

4. check where it is?

```
root@storage:~# lvsdisplay
--- Logical volume ---
LV Path                /dev/ubuntu-vg/root
LV Name                root
VG Name                ubuntu-vg
LV UUID                NA7DgH-V0Sv-cH8E-wvej-aJmP-6EB0-joX00C
LV Write Access        read/write
LV Creation host, time ubuntu, 2017-08-23 16:30:36 +0800
LV Status              available
# open                 1
LV Size                45.52 GiB
Current LE             11653
Segments               1
Allocation             inherit
Read ahead sectors     auto
- currently set to    256
Block device           252:0

--- Logical volume ---
LV Path                /dev/ubuntu-vg/swap_1
LV Name                swap_1
VG Name                ubuntu-vg
LV UUID                Vtixi8-qKcP-f1LH-bHqM-E73h-NN7z-eSD2zk
LV Write Access        read/write
LV Creation host, time ubuntu, 2017-08-23 16:30:36 +0800
LV Status              available
# open                 2
LV Size                4.00 GiB
Current LE             1024
Segments               1
Allocation             inherit
Read ahead sectors     auto
- currently set to    256
Block device           252:1

--- Logical volume ---
LV Path                /dev/cinder-volumes/volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222
LV Name                volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222
VG Name                cinder-volumes
LV UUID                6jbPGA-i3Eo-O4ng-8Mf3-IoeF-9WF7-g1DGEA
LV Write Access        read/write
LV Creation host, time storage, 2017-08-24 22:38:28 +0800
LV Status              available
# open                 0
LV Size                1.00 GiB
Current LE             256
Segments               1
Allocation             inherit
Read ahead sectors     auto
- currently set to    256
Block device           252:2
```

Attach the volume to an instance

1. Attach a volume to an instance:

```
$ openstack server add volume INSTANCE_NAME VOLUME_NAME
```

Replace `INSTANCE_NAME` with the name of the instance and `VOLUME_NAME` with the name of the volume you want to attach to it.

Example

Attach the `volume1` volume to the `provider-instance` instance:

```
$ openstack server add volume provider-instance volume1
```

This command provides no output.

2. List volumes:

```
root@controller:~# openstack volume list
```

```
+-----+-----+-----+-----+
| ID                                     | Display Name | Status | Size | Attached to |
+-----+-----+-----+-----+
| 81ffed40-ed71-495d-bfa9-8fb8c72cf222 | volume1      | in-use | 1    | Attached to |
| provider-instance on /dev/vdb        |              |        |      |              |
+-----+-----+-----+-----+
```

```
root@storage:~# lvdisplay
```

```
--- Logical volume ---
```

```
LV Path                /dev/ubuntu-vg/root
LV Name                 root
VG Name                 ubuntu-vg
LV UUID                 NA7DgH-V0Sv-CH8E-wvej-aJmP-6EBO-joX00C
LV Write Access         read/write
LV Creation host, time  ubuntu, 2017-08-23 16:30:36 +0800
LV Status                available
# open                  1
LV Size                 45.52 GiB
Current LE              11653
Segments                1
Allocation              inherit
Read ahead sectors      auto
- currently set to     256
Block device            252:0
```

```
--- Logical volume ---
```

```
LV Path                /dev/ubuntu-vg/swap_1
LV Name                 swap_1
VG Name                 ubuntu-vg
LV UUID                 Vtixi8-qKcP-f1LH-bHqM-E73h-NN7z-eSD2zk
LV Write Access         read/write
LV Creation host, time  ubuntu, 2017-08-23 16:30:36 +0800
LV Status                available
# open                  2
LV Size                 4.00 GiB
Current LE              1024
Segments                1
Allocation              inherit
Read ahead sectors      auto
- currently set to     256
Block device            252:1
```

```
--- Logical volume ---
```

```
LV Path                /dev/cinder-volumes/volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222
LV Name                 volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222
VG Name                 cinder-volumes
LV UUID                 6jbPGA-i3Eo-O4ng-8Mf3-IoeF-9WF7-g1DGEA
LV Write Access         read/write
```

```

LV Creation host, time storage, 2017-08-24 22:38:28 +0800
LV Status          available
# open             1
LV Size            1.00 GiB
Current LE         256
Segments           1
Allocation          inherit
Read ahead sectors auto
- currently set to 256
Block device       252:2

```

3. Access your instance using SSH or `virsh console` and use the `fdisk` command to verify presence of the volume as the `/dev/vdb` block storage device:

```

$ sudo fdisk -l

Disk /dev/vda: 1073 MB, 1073741824 bytes
255 heads, 63 sectors/track, 130 cylinders, total 2097152 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk identifier: 0x00000000

   Device Boot      Start         End      Blocks   Id  System
/dev/vda1    *        16065      2088449      1036192+   83   Linux

Disk /dev/vdb: 1073 MB, 1073741824 bytes
16 heads, 63 sectors/track, 2080 cylinders, total 2097152 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk identifier: 0x00000000

Disk /dev/vdb doesn't contain a valid partition table

```

4. Check from storage node on iSCSI target point of view, it's found

- o Initiator: `iqn.1993-08.org.debian:01:e7b693dedcab alias: compute`
- o LUN 1: `Backing store path: /dev/cinder-volumes/volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222`

```

root@storage:~# tgtadm --lld iscsi --op show --mode target
Target 1: iqn.2010-10.org.openstack:volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222
  System information:
    Driver: iscsi
    State: ready
  I_T nexus information:
    I_T nexus: 1
      Initiator: iqn.1993-08.org.debian:01:e7b693dedcab alias: compute
      Connection: 0
        IP Address: 192.168.199.20
  LUN information:
    LUN: 0
      Type: controller
      SCSI ID: IET      00010000
      SCSI SN: beaf10
      Size: 0 MB, Block size: 1
      Online: Yes
      Removable media: No
      Prevent removal: No
      Readonly: No
      SWP: No
      Thin-provisioning: No
      Backing store type: null
      Backing store path: None
      Backing store flags:
    LUN: 1
      Type: disk
      SCSI ID: IET      00010001
      SCSI SN: beaf11
      Size: 1074 MB, Block size: 512
      Online: Yes
      Removable media: No
      Prevent removal: No
      Readonly: No
      SWP: No
      Thin-provisioning: No
      Backing store type: rdwr
      Backing store path: /dev/cinder-volumes/volume-81ffed40-ed71-495d-bfa9-
8fb8c72cf222
      Backing store flags:
  Account information:
    7jTnhhxXsVM4BwqxG979
  ACL information:
    ALL

```

5. Checking from compute via `virsh dumpxml <instance-id>`

It's shown the device from initiator point of view:

```
<disk type='block' device='disk'>
  <driver name='qemu' type='raw' cache='none' io='native' />
  <source dev='/dev/disk/by-path/ip-192.168.199.30:3260-iscsi-iqn.2010-
10.org.openstack:volume-81ffed40-ed71-495d-bfa9-8fb8c72cf222-lun-1' />
  <backingStore />
  <target dev='vdb' bus='virtio' />
  <serial>81ffed40-ed71-495d-bfa9-8fb8c72cf222</serial>
  <alias name='virtio-disk1' />
  <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x0' />
</disk>
```